

exploration

November 21, 2025

1 Comparação de Modelos (Exploration)

Notebook para comparar diferentes classificadores no pipeline de Predictive RCA, mantendo boas práticas: divisão estratificada, pré-processamento único, random seed fixo e avaliação consistente em validação e teste.

Notas de boas práticas

- Manter o mesmo pré-processamento e splits para todos os modelos (evita vazamento e garante comparabilidade).
- Fixar `random_state` para reproduzibilidade.
- Avaliar em validação e teste com a mesma métrica principal (AUC-ROC) e threshold padrão 0.5 para classificação.
- Inspecionar relatórios por classe para checar precisão/recall e usar gráficos (ROC, matriz de confusão) para entender a resposta de cada algoritmo.

```
[15]: import sys
import os

sys.path.append(os.path.abspath("../"))
```

```
[16]: from pathlib import Path
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import (
    ConfusionMatrixDisplay,
    classification_report,
    confusion_matrix,
    roc_auc_score,
    roc_curve,
)

from src.config.settings import CONFIG
from src.preprocessing.build_features import build_case_features
from src.preprocessing.split import stratified_case_split
from src.preprocessing.preprocess import (
    build_preprocessor,
    fit_preprocessor_and_transform,
)
```

```

from src.models.lightgbm_model import LightGBMModel
from src.models.random_forest import RandomForestModel
from src.models.logistic_regression import LogisticRegressionModel
from src.models.xgboost_model import XGBoostModel

pd.options.display.float_format = lambda v: f"{v:.4f}"
plt.style.use("ggplot")

```

[17]: # Carregar e preparar dataset em nível de caso

```

log_path = Path("../data/raw/event_log_sintetico_2000_cases.csv")
sla_hours = CONFIG.SLA_HOURS

df_events = pd.read_csv(log_path)
if "timestamp" in df_events.columns:
    df_events["timestamp"] = pd.to_datetime(df_events["timestamp"])

df_cases = build_case_features(df_events, sla_hours=sla_hours)
df_cases.head()

```

[17]:

	case_id	throughput_hours	num_events	num_unique_activities	rework_count	\
0	1	46.3033	4		4	0
1	2	69.1386	5		5	0
2	3	8.9614	6		5	1
3	4	41.3628	7		5	2
4	5	26.1333	6		6	0

	start_activity	end_activity	start_resource	mean_cost	sla_violated
0	Receber Pedido	Finalizar	Ana	87.5350	1
1	Receber Pedido	Finalizar	João	72.2880	1
2	Receber Pedido	Finalizar	Pedro	79.4467	0
3	Receber Pedido	Finalizar	João	75.1771	1
4	Receber Pedido	Finalizar	Marina	79.8367	1

[18]: # Split estratificado e pré-processamento único para todos os modelos

```

splits, numeric_features, categorical_features = stratified_case_split(
    df_cases,
    target_col="sla_violated",
    id_col="case_id",
    random_state=CONFIG.RANDOM_STATE,
)

preprocessor = build_preprocessor(numeric_features, categorical_features)
artifacts, X_train_pre, X_val_pre, X_test_pre = fit_preprocessor_and_transform(
    preprocessor, splits, numeric_features, categorical_features
)

y_train = splits.y_train.to_numpy()

```

```

y_val = splits.y_val.to_numpy()
y_test = splits.y_test.to_numpy()

numeric_features, categorical_features

```

```
[18]: (['throughput_hours',
      'num_events',
      'num_unique_activities',
      'rework_count',
      'mean_cost'],
      ['start_activity', 'end_activity', 'start_resource'])
```

```
[19]: from typing import Any, Dict, List, Tuple
```

```

def summarize_split(y_true, y_proba, threshold: float = 0.5):
    y_pred = (y_proba >= threshold).astype(int)
    report_dict = classification_report(
        y_true, y_pred, output_dict=True, zero_division=0
    )
    return {
        "auc": roc_auc_score(y_true, y_proba),
        "precision": report_dict["1"]["precision"],
        "recall": report_dict["1"]["recall"],
        "f1": report_dict["1"]["f1-score"],
        "support": int(report_dict["1"]["support"]),
        "report_text": classification_report(y_true, y_pred, zero_division=0),
        "confusion": confusion_matrix(y_true, y_pred),
        "pred": y_pred,
        "proba": y_proba,
    }

def train_and_eval(
    model_class, params: Dict[str, Any] | None = None, name: str | None = None
):
    params = params or {}
    model = model_class(**params)
    model.train(X_train_pre, y_train)

    val = summarize_split(y_val, model.predict_proba(X_val_pre)[:, 1])
    test = summarize_split(y_test, model.predict_proba(X_test_pre)[:, 1])

    return {
        "model": name or model_class.__name__,
        "val_auc": val["auc"],
        "test_auc": test["auc"],
    }

```

```

    "val_precision": val["precision"],
    "test_precision": test["precision"],
    "val_recall": val["recall"],
    "test_recall": test["recall"],
    "val_f1": val["f1"],
    "test_f1": test["f1"],
    "val_report": val["report_text"],
    "test_report": test["report_text"],
    "val_confusion": val["confusion"],
    "test_confusion": test["confusion"],
    "val_pred": val["pred"],
    "test_pred": test["pred"],
    "val_proba": val["proba"],
    "test_proba": test["proba"],
}

experiments: List[Tuple[str, Any, Dict[str, Any]]] = [
(
    "LightGBM",
    LightGBMModel,
    {
        "objective": "binary",
        "n_estimators": 300,
        "learning_rate": 0.05,
        "num_leaves": 64,
        "max_depth": -1,
        "feature_fraction": 0.8,
        "random_state": CONFIG.RANDOM_STATE,
        "n_jobs": -1,
    },
),
(
    "RandomForest",
    RandomForestModel,
    {
        "n_estimators": 400,
        "max_depth": 12,
        "min_samples_leaf": 2,
        "random_state": CONFIG.RANDOM_STATE,
        "n_jobs": -1,
    },
),
(
    "LogisticRegression",
    LogisticRegressionModel,
    {

```

```

        "max_iter": 1000,
        "class_weight": "balanced",
        "n_jobs": -1,
    },
),
(
    "XGBoost",
    XGBoostModel,
{
    "n_estimators": 300,
    "learning_rate": 0.05,
    "max_depth": 6,
    "subsample": 0.8,
    "colsample_bytree": 0.8,
    "random_state": CONFIG.RANDOM_STATE,
    "n_jobs": -1,
},
),
]

```

```
[20]: results = []
for name, cls, params in experiments:
    res = train_and_eval(cls, params, name=name)
    results.append(res)

metrics_df = (
    pd.DataFrame(results)[
        [
            "model",
            "val_auc",
            "test_auc",
            "val_precision",
            "test_precision",
            "val_recall",
            "test_recall",
            "val_f1",
            "test_f1",
        ]
    ]
    .sort_values("val_auc", ascending=False)
    .reset_index(drop=True)
)
metrics_df
```

[LightGBM] [Warning] feature_fraction is set=0.8, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.8
[LightGBM] [Warning] feature_fraction is set=0.8, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.8


```

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] feature_fraction is set=0.8, colsample_bytree=1.0 will be
ignored. Current value: feature_fraction=0.8
[LightGBM] [Warning] feature_fraction is set=0.8, colsample_bytree=1.0 will be
ignored. Current value: feature_fraction=0.8

/Users/pedroeckel/.local/share/virtualenvs/predictive-
rca-9dRCaVbd/lib/python3.13/site-packages/sklearn/utils/validation.py:2749:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
    warnings.warn(
/Users/pedroeckel/.local/share/virtualenvs/predictive-
rca-9dRCaVbd/lib/python3.13/site-packages/sklearn/utils/validation.py:2749:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
    warnings.warn(
/Users/pedroeckel/.local/share/virtualenvs/predictive-
rca-9dRCaVbd/lib/python3.13/site-packages/xgboost/training.py:199: UserWarning:
[01:30:14] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

        bst.update(dtrain, iteration=i, fobj=obj)

```


```

[20]:      model  val_auc  test_auc  val_precision  test_precision \
0       LightGBM  1.0000  1.0000      1.0000      1.0000
1   RandomForest  1.0000  1.0000      1.0000      1.0000
2 LogisticRegression  1.0000  1.0000      1.0000      1.0000
3       XGBoost  1.0000  1.0000      1.0000      1.0000

```

	val_recall	test_recall	val_f1	test_f1
0	1.0000	0.9969	1.0000	0.9984
1	1.0000	1.0000	1.0000	1.0000
2	0.9969	0.9907	0.9984	0.9953
3	1.0000	1.0000	1.0000	1.0000

```
[21]: ax = metrics_df.set_index("model")[["val_auc", "test_auc"]].plot.bar(
    figsize=(8, 4), ylim=(0, 1)
)
ax.set_ylabel("AUC-ROC")
ax.set_title("AUC-ROC por modelo (validação vs. teste)")
ax.legend(title="Split")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```



```
[22]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))
for res in results:
    fpr_val, tpr_val, _ = roc_curve(y_val, res["val_proba"])
    fpr_test, tpr_test, _ = roc_curve(y_test, res["test_proba"])
    axes[0].plot(
        fpr_val,
        tpr_val,
        lw=2,
        label=f"{res['model']} (AUC={res['val_auc']:.3f})",
    )
    axes[1].plot(
```


```

        fpr_test,
        tpr_test,
        lw=2,
        label=f"{res['model']} (AUC={res['test_auc']:.3f})",
    )

for ax, title in zip(axes, ["Validação", "Teste"]):
    ax.plot([0, 1], [0, 1], color="gray", linestyle="--", alpha=0.5)
    ax.set_xlim(0, 1)
    ax.set_ylim(0, 1.02)
    ax.set_title(f"Curvas ROC - {title}")
    ax.set_xlabel("Falsa positividade")
    ax.set_ylabel("Verdadeira positividade")
    ax.legend()

plt.tight_layout()
plt.show()

```




```

[23]: fig, axes = plt.subplots(len(results), 2, figsize=(10, 3 * len(results)))
axes = axes if len(results) > 1 else [axes]
for i, res in enumerate(results):
    row_axes = axes[i]
    for ax, (split_label, cm) in zip(
        row_axes,
        [("Validação", res["val_confusion"]), ("Teste", res["test_confusion"])],
    ):
        disp = ConfusionMatrixDisplay(cm, display_labels=["Dentro SLA", "Violou SLA"])
        disp.plot(ax=ax, colorbar=False, cmap="Blues")
        ax.set_title(f"{res['model']} - {split_label}")

```

```
plt.tight_layout()  
plt.show()
```



```
[24]: class_map = {"0": "Dentro do SLA", "1": "Violou SLA"}
report_rows = []
for res in results:
    for split_label, y_true, y_pred in [
        ("Validação", y_val, res["val_pred"]),
        ("Teste", y_test, res["test_pred"]),
    ]:
        report_dict = classification_report(
            y_true, y_pred, output_dict=True, zero_division=0
        )
        for cls in ("0", "1"):
            stats = report_dict[cls]
            report_rows.append(
                {
                    "modelo": res["model"],
                    "split": split_label,
                    "classe": class_map[cls],
                    "precisao": stats["precision"],
                    "recall": stats["recall"],
                    "f1": stats["f1-score"],
                    "suporte": int(stats["support"]),
                }
            )
    )

per_class_reports = pd.DataFrame(report_rows)
per_class_reports
```

	modelo	split	classe	precisao	recall	f1	\
0	LightGBM	Validação	Dentro do SLA	1.0000	1.0000	1.0000	
1	LightGBM	Validação	Violou SLA	1.0000	1.0000	1.0000	
2	LightGBM	Teste	Dentro do SLA	0.9875	1.0000	0.9937	
3	LightGBM	Teste	Violou SLA	1.0000	0.9969	0.9984	
4	RandomForest	Validação	Dentro do SLA	1.0000	1.0000	1.0000	
5	RandomForest	Validação	Violou SLA	1.0000	1.0000	1.0000	
6	RandomForest	Teste	Dentro do SLA	1.0000	1.0000	1.0000	
7	RandomForest	Teste	Violou SLA	1.0000	1.0000	1.0000	
8	LogisticRegression	Validação	Dentro do SLA	0.9875	1.0000	0.9937	
9	LogisticRegression	Validação	Violou SLA	1.0000	0.9969	0.9984	
10	LogisticRegression	Teste	Dentro do SLA	0.9634	1.0000	0.9814	
11	LogisticRegression	Teste	Violou SLA	1.0000	0.9907	0.9953	
12	XGBoost	Validação	Dentro do SLA	1.0000	1.0000	1.0000	
13	XGBoost	Validação	Violou SLA	1.0000	1.0000	1.0000	
14	XGBoost	Teste	Dentro do SLA	1.0000	1.0000	1.0000	
15	XGBoost	Teste	Violou SLA	1.0000	1.0000	1.0000	
	suporte						
0	79						

```
1      321
2      79
3      321
4      79
5      321
6      79
7      321
8      79
9      321
10     79
11     321
12     79
13     321
14     79
15     321
```

```
[25]: reports_text = {
    res["model"] : {"validacao": res["val_report"], "teste": res["test_report"]}
    for res in results
}
reports_text
```

```
[25]: {'LightGBM': {'validacao': 'precision      recall   f1-score
support\n\n          0      1.00      1.00      1.00      79\n          1
1.00      1.00      1.00      321\n\n      accuracy
1.00      400\nn      macro avg      1.00      1.00      1.00      400\nnweighted
avg      1.00      1.00      1.00      400\n', 'teste': 'precision      recall   f1-score
support\n\n          0      1.00      0.99      79\n          1      1.00      1.00
1.00      321\n\n      accuracy
1.00      400\nn      macro avg      0.99      1.00      1.00      400\nnweighted
avg      1.00      1.00      1.00      400\n', 'RandomForest': {'validacao': 'precision      recall   f1-score
support\n\n          0      1.00      1.00      1.00      79\n          1
1.00      1.00      1.00      321\n\n      accuracy
1.00      400\nn      macro avg      1.00      1.00      1.00      400\nnweighted
avg      1.00      1.00      1.00      400\n', 'teste': 'precision      recall   f1-score
support\n\n          0      1.00      1.00      79\n          1      1.00      1.00
1.00      321\n\n      accuracy
1.00      400\nn      macro avg      1.00      1.00      1.00      400\nnweighted
avg      1.00      1.00      1.00      400\n', 'LogisticRegression': {'validacao': 'precision      recall
f1-score      support\n\n          0      0.99      1.00      0.99      79\n          1
1.00      1.00      1.00      1.00      321\n\n      accuracy
1.00      400\nn      macro avg      0.99      1.00      1.00      400\nnweighted
avg      1.00      1.00      1.00      400\n', 'teste': ''}}
```

```

'teste': 'precision      recall   f1-score   support\n\n
0       0.96      1.00      0.98      79\n      1       1.00      0.99\n1.00    321\naccuracy\nmacro avg      0.98      1.00      0.99      400\nweighted avg      0.99\n0.99      0.99      400\n}',

'XGBoost': {'validacao': 'precision      recall   f1-score\nsupport\n\n      0       1.00      1.00      1.00      79\n      1\n1.00    1.00      1.00      321\naccuracy\n1.00    400\nmacro avg      1.00      1.00      1.00      400\nweighted\navg      1.00      1.00      1.00      400\n',\n'teste': 'precision      recall   f1-score   support\n\n0       1.00      1.00      1.00      79\n      1       1.00      1.00\n1.00    321\naccuracy\nmacro avg      1.00      1.00      1.00      400\nweighted avg      1.00\n1.00    1.00      400\n'}\n

```